

# Assignment 3: Sentiment Classification

## Learning Curves

### 1. Random Embeddings



Sentiment Classification Accuracy of 52.20%

Confusion matrix:

[[518. 393.]

[477. 432.]]

## 2. GloVe Embeddings



Sentiment Classification Accuracy of 84.12%

Confusion matrix:

[[731. 180.]

[109. 800.]]

### 3. ELMO Embeddings



Sentiment Classification Accuracy of 89.01%

Confusion matrix:

[[782. 129.]

[ 71. 838.]]

## 4. ELMO + GloVe Embeddings



Sentiment Classification Accuracy of 89.62%

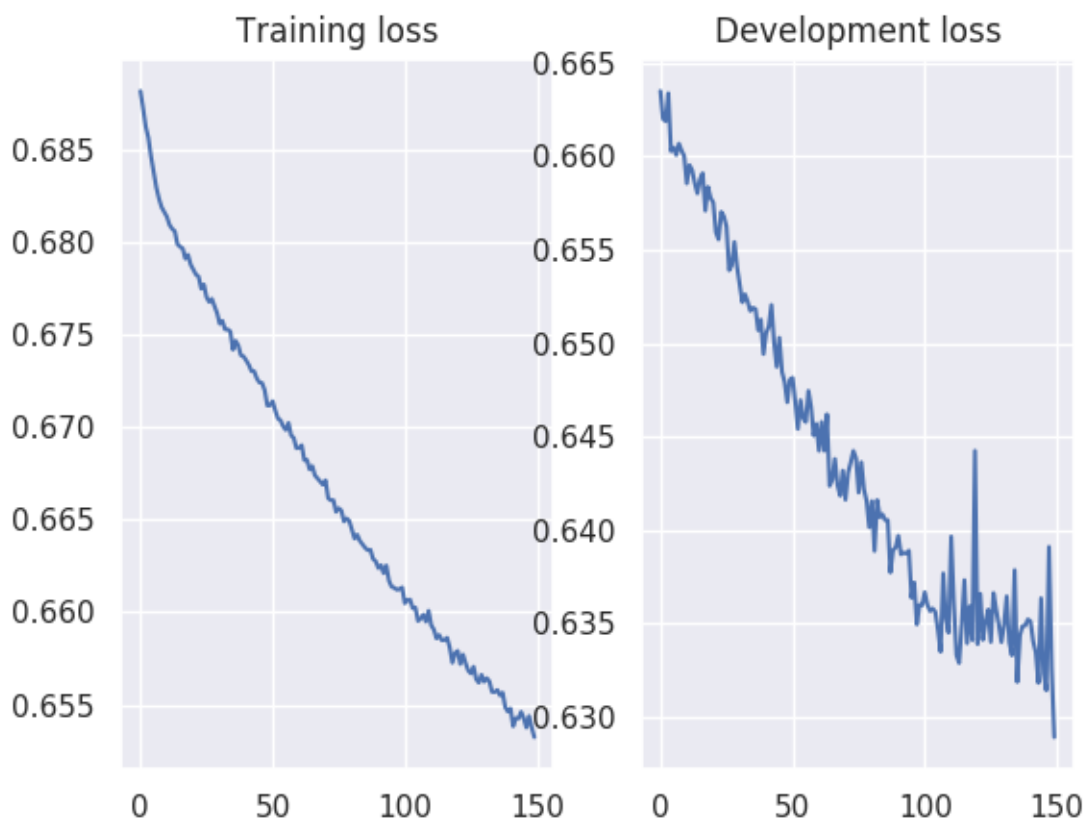
Confusion matrix:

```
tensor([[791., 120.],  
        [ 69., 840.]])
```

## Hyperparameter Experiments

**Experiment 1: Does training the Sentiment Network with random embeddings over more epochs rival ELMO/Glove Embeddings?**

Motivation: Training with random embeddings is a lot faster than training with Elmo or even Glove. If we look at the training loss curve from training with random embeddings, it has not plateaued by 10 epochs. I suspect that if we just train with random embeddings for more epochs, its performance could be more comparable to Elmo or Glove.



Results: As you can see above, I trained with random embeddings for 150 episodes. Training time was similar to Glove embeddings and lower than Elmo embeddings.

Sentiment Classification Accuracy of 59.12%

Confusion matrix:

```
tensor([[538., 373.],
        [371., 538.]])
```

The classification accuracy went up from 52% (10 epochs) to 59% (150 epochs). However, the classification accuracy didn't come close to the >80% accuracy when using Glove or Elmo.

***Experiment 2: Impact of increasing the RNN hidden size on Glove vs Elmo***

The results in the previous section were obtained with the given hyperparameter setting of the RNN hidden size of 1024. I changed that hidden size to 2048 and observed the impact on validation accuracy for glove and elmo embeddings.

Results:

Glove Embeddings with 2046 hidden size:

Sentiment Classification Accuracy of 82.14%

Confusion matrix:

```
tensor([[702., 209.],  
        [116., 793.]])
```

Elmo Embeddings with 2048 embedding size:

**Sentiment Classification Accuracy of 90.27%**

Confusion matrix:

```
tensor([[828., 83.],  
        [ 94., 815.]])
```

Thus we can see that increasing the layer size had a big impact on sentiment classification with ELMO.

## Overfitting

For all embedding methods, my loss curve on the validation data was increasing with the number of epochs – suggesting severe overfitting of the model to the training data. To mitigate overfitting, I introduced dropout (with probability of dropping a neuron set to 0.5) after the non-linear activations of the classification module in SentimentNetwork.

I reduced the capacity of my classification module by using just one linear layer with 2 intermediate nodes. While this resulted in the dev loss to fall with more epochs, it resulted in my classifier always predicting the same class. As a result, I decided that I was going to have to overfit to the training data in order to get a better classification accuracy.

I also found that if I trained for more epochs (for random embeddings), I was no longer overfitting and the dev loss started to go down.